



Politecnico
di Torino



Machine Learning for Networking

Pandas

DataBase and Data Mining Group

Andrea Pasini
Flavio Giobergia
Elena Baralis
Gabriele Ciravegna



Introduction to Pandas

- Pandas
 - Provides useful data structures (Series and DataFrames) and data analysis tools
 - Based on **Numpy** arrays
 - Data analysis tools:
 - Managing **tables** and **series**
 - data selection
 - grouping, pivoting
 - **Statistics** on data
 - Managing **missing data** (extra slides)



- **Series:** 1-Dimensional sequence of homogeneous elements (“values”)
- Elements are associated to an explicit **index**
 - index elements can be either strings or integers
- Examples:

{	index	1	2	3		
	values	0.3	0.5	0.8		

{	index	'3-July'	'4-July'	'5-July'		
	values	0.3	0.5	0.8		



■ Creation from list

- When not specified, index is set automatically with a progressive number



In [1]:

```
import pandas as pd  
s1 = pd.Series([2.0, 3.1, 4.5])  
print(s1)
```

Out[1]:

```
0    2.0  
1    3.1  
2    4.5
```



- **Creation** from list, specifying index



In [1]: `pd.Series([2.0, 3.1, 4.5], index=['mon', 'tue', 'wed'])`

Out[1]:

'mon'	2.0
'tue'	3.1
'wed'	4.5

Strings as
index
elements



- **Creation** from dictionary
 - The keys define the index
 - The values of the dictionary the values of the series

```
In [1]: pd.Series({'c':2.0, 'b':3.1, 'a':4.5})
```

```
Out[1]:
```

'c'	2.0
'b'	3.1
'a'	4.5



- Obtaining **values** and **index** from a Series



```
In [1]: s1 = pd.Series([2.0, 3.1, 4.5], index=['mon', 'tue', 'wed'])
        print(s1.values)
        print(s1.index)
```

```
Out[1]: [2.0, 3.1, 4.5] # Numpy array
        Index(['mon', 'tue', 'wed'], dtype='object')
```

- The **Index** is a custom Python object defined in Pandas



- Accessing Series elements
- **Access by Index**
 - **Explicit:** the one specified while creating a Series
 - Series.**loc** attribute
 - **Implicit:** number associated to the element order (similarly to List or Numpy arrays accessing)
 - Series.**iloc** attribute
- In both cases you use `s.loc[index]` or `s.iloc[index]`



■ Accessing Series elements



```
In [1]: s1 = pd.Series([2.0, 3.1, 4.5], index=['a', 'b', 'c'])
print(s1.loc['a'])           # With explicit index
print(s1.iloc[0])           # With implicit index
s1.loc['b'] = 10             # Both Allows editing values
print(f"Series:\n{s1}")
```

```
Out[1]: 2.0
2.0 # We return the same value
Series:
'a'      2.0
'b'      10
'c'      4.5
```



■ Accessing Series elements: **slicing**



```
In [1]: s1 = pd.Series([2.0, 3.1, 4.5], index=['a', 'b', 'c'])  
print(s1.loc['b':'c']) # explicit index (stop element included)  
print(s1.iloc[1:3])   # implicit index (stop element excluded)
```

```
Out[1]:  
b 3.1  
c 4.5  
  
b 3.1  
c 4.5
```

Very
dangerous for
numerical
indexes!



■ Accessing Series elements: **masking**



In [1]:

```
s1 = pd.Series([2.0, 3.1, 4.5], index=['a', 'b', 'c'])
print(s1.loc[(s1>2) & (s1<10)])

# directly indexing the series also works
# print(s1[(s1>2)& (s1<10)])
# but try to avoid it! It creates problems with dataframes
```

Out[1]:

```
b 3.1
c 4.5
```



- Accessing Series elements: **fancy indexing**



```
In [1]: s1 = pd.Series([2.0, 3.1, 4.5], index=['a', 'b', 'c'])  
print(s1.loc[['a', 'c']])  
print(s1.iloc[[0, 2]])
```

```
Out[1]:  
  
a 2.0  
c 4.5  
  
a 2.0  
c 4.5
```



- **DataFrame**: 2-Dimensional array

- Like a table where:
 - **Columns are Series** objects
 - Each column has a **name**
 - Columns share the **same index**

- **Example**:

Index	'Price'	'Quantity'	'Liters'
'Water'	1.0	5	1.5
'Beer'	1.4	10	0.3
'Wine'	5.0	8	1



■ Creation from Series

- Use a **dictionary** to set column names



Series should share the same index

```
In [1]: price = pd.Series([1.0, 1.4, 5], index=['a', 'b', 'c'])
quantity = pd.Series([5, 10, 8], index=['a', 'b', 'c'])
liters = pd.Series([1.5, 0.3, 1], index=['a', 'b', 'c'])
df = pd.DataFrame({'Price':price, 'Quantity':quantity,
                  'Liters':liters})

print(df)
```

```
Out[1]:
```

	Price	Quantity	Liters
a	1.0	5	1.5
b	1.4	10	0.3
c	5.0	8	1.0



- **Creation** from dictionary of key-list pairs
 - **Each value (list)** is associated to a **column**
 - Column name given by the key
 - All lists should have the same length
 - **Index** is automatically set to a progressive number
 - Unless explicitly passed as parameter (index=...)
- **Example:**

```
In [1]: dct = { "c1": [0, 1, 2], "c2": [0, 2, 4] }  
        df = pd.DataFrame(dct)  
        print(df)
```

```
Out[1]:
```

	c1	c2
0	0	0
1	1	2
2	2	4



- **Creation** from list of dictionaries
 - **Each dictionary** is associated to a **row**
 - All dictionary should have the same keys
 - **Index** is automatically set to a progressive number
 - Unless explicitly passed as parameter (index=...)

- **Example:**

```
In [1]: dic_list = [{'c1':i, 'c2':2*i} for i in range(3)]
df = pd.DataFrame(dic_list)
print(df)
```

```
Out[1]:
```

	c1	c2
0	0	0
1	1	2
2	2	4



- **Creation** from 2D Numpy array
- **Example:**



```
In [1]: arr = np.arange(6).reshape((3,2))
df = pd.DataFrame(arr, columns=['c1', 'c2'],
                  index=['a', 'b', 'c'])
print(df)
```

```
Out[1]:
```

	c1	c2
a	0	1
b	2	3
c	4	5



- `df.columns`, `df.index`: to obtain **column names** and **index** of a DataFrame

Index	Price	Quantity	Liters
a	1.0	5	1.5
b	1.4	10	0.3
c	5.0	8	1

```
In [2]: print(df.columns) # Index object with column names  
        print(df.index)  # Index object
```

```
Out[2]: Index(['Price', 'Quantity', 'Liters'], dtype='object')  
        Index(['a', 'b', 'c'], dtype='object')
```



- `df.values`: to **access** DataFrame data
 - Get a 2D Numpy array

Index	Price	Quantity	Liters
a	1.0	5	1.5
b	1.4	10	0.3
c	5.0	8	1

In [2]: `print(df.values) # Numpy array with data`

Out[2]: `array([[1.0, 5.0, 1.5],
 [1.4, 10.0, 0.3],
 [5.0, 8.0, 1.0]])`



- **Accessing DataFrames**
 - Access a DataFrame column
 - Access rows and columns with indexing
 - **df.loc**
 - **Explicit** index
 - Slicing, masking, fancy indexing
 - **df.iloc**
 - **Implicit** index
- Whether a **copy** or **view** will be returned it depends on the context
 - Usually it is difficult to make assumptions
 - Use inline assignment with explicit indexing `.loc`
 - https://pandas-docs.github.io/pandas-docs-travis/user_guide/indexing.html



- **Accessing DataFrame columns**
 - By directly indexing the column name
 - Returns a **Series** with column data



This is why we should **avoid** using direct indexing for **row** indexing

Index	Price	Quantity	Liters
a	1.0	5	1.5
b	1.4	10	0.3
c	5.0	8	1

```
In [1]: df['Quantity']  
# df.loc['Quantity'] does not work
```

```
Out[1]:  
a      5  
b     10  
c      8
```



- Direct indexing:
 - Access single DataFrame **row** by index
 - **loc** (explicit), **iloc** (implicit)
 - Return a **Series** with an element for each column



```
In [1]: print(df.loc['a'])           # Get the first row (explicit)
        print(df.iloc[0])           # Get the first row (implicit)
```

```
Out[1]: Price      1.0
         Quantity  5.0
         Liters    1.5

         Price      1.0
         Quantity  5.0
         Liters    1.5
```



■ Accessing DataFrames with **slicing**

- Allows selecting rows and columns

- As in numpy **columns** are the last dimension, **rows** the second to last



In [1]:

```
print(df.loc['b':'c', 'Quantity':'Liters'])
```

Out[1]:

	Quantity	Liters
b	10	0.3
c	8	1



■ Accessing DataFrames with **masking**

- Select rows based on a condition



Index	Price	Quantity	Liters
a	1.0	5	1.5
b	1.4	10	0.3
c	5.0	8	1

```
In [1]: mask = (df['Quantity']<10) & (df['Liters']>1)
df.loc[mask, 'Quantity':] # Use masking and slicing
```

```
Out[1]:
```

	Quantity	Liters
a	5	1.5



■ Accessing DataFrames with fancy indexing



- To select **columns**...

Index	Price	Quantity	Liters
a	1.0	5	1.5
b	1.4	10	0.3
c	5.0	8	1

```
In [1]: mask = (df['Quantity']<10) & (df['Liters']>1)
df.loc[mask, ['Price','Liters']] # Use masking and fancy
```

```
Out[1]:
```

	Price	Liters
a	1.0	1.5



■ Accessing DataFrames with **fancy indexing**

- To select **rows** and **columns**...



Index	Price	Quantity	Liters
a	1.0	5	1.5
b	1.4	10	0.3
c	5.0	8	1

In [1]: `df.loc[['a', 'c'], ['Price', 'Liters']]`

Out[1]:

	Price	Liters
a	1.0	1.5
c	5.0	1.0



■ Assign value to selected items

```
In [1]: df.loc[['a', 'c'], ['Price', 'Liters']] = 0
```

With inline assignment we are not worry about views or copy

Index	Price	Quantity	Liters
a	0.0	5	0.0
b	1.4	10	0.3
c	0.0	8	0.0



- **Add new column** to DataFrame
 - DataFrame is modified **inplace**

Index	Price	Quantity	Liters		Index	Price	Quantity	Liters	Available
a	0.0	5	0.0		a	1.0	5	1.5	True
b	1.4	10	0.3	→	b	1.4	10	0.3	False
c	0.0	8	0.0		c	5.0	8	1	True

In [1]:

```
df['Available'] = pd.Series([True, False, True],  
                             index=['a', 'b', 'c'])
```

- If the DataFrame already has a column with the specified name, then this is **replaced**



- **Add new column** to DataFrame
 - It is also possible to assign directly a **list**

Index	Price	Quantity	Liters
a	0.0	5	0.0
b	1.4	10	0.3
c	0.0	8	0.0



Index	Price	Quantity	Liters	Available
a	1.0	5	1.5	True
b	1.4	10	0.3	False
c	5.0	8	1	True

In [1]: `df['Available'] = [True, False, True]`



■ Drop column(s)

- Returns a **copy** of the updated DataFrame
 - Unless inplace=True, in which case the original DataFrame is modified
 - This applies to many pandas methods -- always check the documentation!

Index	Price	Quantity	Liters	Available
a	1.0	5	1.5	True
b	1.4	10	0.3	False
c	5.0	8	1	True

In [1]:

```
df = df.drop(columns=['Quantity', 'Liters'])
```



■ Rename column(s)

- Use a **dictionary** which maps old names with new names
- Returns a **copy** of the updated DataFrame

Index	Price	Quantity	Liters	Available
a	1.0	5	1.5	True
b	1.4	10	0.3	False
c	5.0	8	1	True



Index	Price	nItems	[L]	Available
a	1.0	5	1.5	True
b	1.4	10	0.3	False
c	5.0	8	1	True


In [1]:

```
df = df.rename(columns={'Quantity': 'nItems',  
                        'Liters': '[L]'})
```



Notebook Examples

- **3.1 Pandas Series and DataFrames.ipynb**





- Unary operations on Series and DataFrames
 - exponentiation, logarithms, ...
- Operations between Series and DataFrames
 - Operations are performed **element-wise**, being aware of their **indices/columns**
- Aggregations (min, max, std, ...)



- Unary operations on Series and DataFrames
 - They work with any **Numpy** ufunc
 - The operation is applied to each element of the Series/DataFrame
 - The result is still a Series/Dataframe!
- Examples:
 - `res = my_series/4 + 1`
 - `res = np.abs(my_series)`
 - `res = np.exp(my_dataframe)`
 - `res = np.sin(my_series/4)`
 - ...



- Operations between Series (+, -, *, /)
 - Applied element-wise after **aligning indices**
 - Index elements which do not match are set to **NaN** (Not a Number)

- Example:

- `res = my_series1 + my_series2`

Index	
b	3
a	1
c	10

my_series1

Index	
a	1
b	3
d	30

my_series2

Index	
a	2
b	6
c	NaN
d	NaN

res

After index alignment
index in the result is **sorted**





- Operations between DataFrames
 - Applied element-wise after **aligning indices** and **columns**
 - Example (align **index**):
 - `res = my_dataframe1 + my_dataframe2`

Index	Total	Quantity
b	3	4
a	1	2
c	10	20

my_dataframe1

Index	Total	Quantity
a	1	2
b	3	4
d	30	40

my_dataframe2

Index in the result
is **sorted**

Index	Total	Quantity
a	2	4
b	6	8
c	NaN	NaN
d	NaN	NaN

res



■ Operations between DataFrames

■ Example (align **columns**)

■ `res = my_dataframe1 + my_dataframe2`

Columns in the result are **sorted**

Index	Total	Quantity
a	1	2
b	3	4
c	5	6

my_dataframe1

Index	Total	Price
a	1	2
b	3	4
c	5	6

my_dataframe2

Index	Price	Quantity	Total
a	NaN	NaN	2
b	NaN	NaN	6
c	NaN	NaN	10

res



- Operations between DataFrames and Series
 - The operation is applied between the Series and each **row** of the DataFrame
 - Follows **broadcasting** rules
 - Example:
 - `res = my_dataframe1 + my_series1`

Index	Total	Quantity
a	1	2
b	3	4
c	5	6

my_dataframe1

Index	
Total	1
Quantity	2

my_series1

Index	Total	Quantity
a	2	4
b	4	6
c	6	8

res



- Pandas Series and DataFrames allow performing aggregations
 - mean, std, min, max, sum
- Examples

```
In [1]: my_series.mean() # Return the mean of Series elements
```

- For DataFrames, aggregate functions are applied **along the column** (i.e. between the rows) and they return a Series long as the number of columns

```
In [1]: my_df.mean() # Return a Series
```



- Example of **aggregations** with DataFrames:
z-score normalization

In [1]:

```
mean_series = df.mean()  
std_series = df.std()  
df_norm = (df-mean_series)/std_series
```

Index	Total	Quantity
a	1	2
b	3	4
c	5	6

my_dataframe1

Index	
Total	3.0
Quantity	4.0

mean_series

Index	
Total	2.0
Quantity	2.0

std_series



Combining Pandas objects

- Pandas provides 2 methods for combining Series and DataFrames
 - `concat()`
 - Concatenate a sequence of Series/DataFrames
 - `append()`
 - Append a Series/DataFrame to the specified object
- 1 method for combining dataframes following relational algebra:
 - `Merge()`
 - Combine two dataframes along certain columns



■ Concatenating 2 Series

- Index is preserved, even if **duplicated**
 - There is nothing that prevents duplicate indices in pandas!

In [1]:

```
s1 = pd.Series(['a', 'b'], index=[1,2])  
s2 = pd.Series(['c', 'd'], index=[1,2])  
pd.concat((s1, s2))
```

Out[1]:

```
1    a  
2    b  
1    c  
2    d  
dtype=object
```



- Concatenating 2 Series
 - To avoid duplicate indexes, use **ignore_index**

In [1]:

```
s1 = pd.Series(['a', 'b'], index=[1,2])  
s2 = pd.Series(['c', 'd'], index=[1,2])  
pd.concat((s1, s2), ignore_index=True)
```

Out[1]:

```
0    a  
1    b  
2    c  
3    d  
dtype=object
```



- Concatenating 2 DataFrames
 - Concatenate **vertically** by default

In [1]:

```
pd.concat((df1, df2))
```

Index	Total	Quantity
a	1	2
b	3	4

Index	Total	Quantity
c	5	6
d	7	8



Index	Total	Quantity
a	1	2
b	3	4
c	5	6
d	7	8



Combining Pandas objects

- Concatenating 2 DataFrames with different columns is possible in Pandas
 - Missing columns are filled with NaN

In [1]:

```
pd.concat((df1, df2))
```

Index	Total	Quantity
a	1	2
b	3	4



Index	Total	Quantity	Liters
c	5	6	1
d	7	8	2

Index	Total	Quantity	Liters
a	1	2	NaN
b	3	4	NaN
c	5	6	1.0
d	7	8	2.0



- The **append()** method is a shortcut for concatenating DataFrames
 - Returns the result of the concatenation

```
In [1]: df_concat = df1.append(df2)
```

is equivalent to:

```
In [1]: df_concat = pd.concat((df1, df2))
```



- Joining DataFrames with relational algebra: **merge()**
 - Merge on:
 - The column(s) with same name in the two DFs, by default
 - Specific columns, by specifying `on=columns`
 - `left_on` and `right_on` may also be used
 - The indices, if `left_index/right_index` are True
 - This preserves the indices (discarded otherwise)
 - Depending on the DataFrames, a **one-to-one**, **many-to-one** or **many-to-many** join can be performed
 - `validate='1:1' | '1:m' | 'm:1' | 'm:m'` to enforce the specific merge

```
In [1]: joined_df = pd.merge(df1, df2)
```



Combining Pandas objects

■ Examples (1)

`pd.merge(df1, df2)` → merge on columns in common, ["k1"]

Index	k1	c2
i1	0	a
i2	1	b

Index	k1	c3
i1	1	b1
i2	0	a1

Index	k1	c2	c3
0	0	a	a1
1	1	b	b1

`pd.merge(df1, df2, right_index=True, left_index=True)` → merge on index

Index	k1	c2
i1	0	a
i2	1	b
i3	0	c
i4	1	d

Index	k1	c3
i1	1	b1
i2	0	a1

Index	k1_x	c2	k1_y	c3
i1	0	a	1	b1
i2	1	b	0	a1



Combining Pandas objects

■ Examples (2)

`pd.merge(df1, df2)` ➔ performs a one-to-one merge

Index	k1	c2	Index	k1	c3	Index	k1	c2	c3
i1	0	a	i1	1	b1	0	0	a	a1
i2	1	b	i2	0	a1	1	1	b	b1

`pd.merge(df1, df2)` ➔ performs a many-to-one merge

Index	k1	c2	Index	k1	c3	Index	k1	c2	c3
i1	0	a	i1	1	b1	0	0	a	a1
i2	1	b	i2	0	a1	1	0	c	a1
i3	0	c				2	1	b	b1
i4	1	d				3	1	d	b1



- Pandas provides the equivalent of the SQL **group by** statement: `df.groupby(...)`
- It allows the following operations:
 - **Iterating** on groups
 - **Aggregating** the values of each group (mean, min, max, ...)
 - **Filtering** groups according to a condition



- **Applying** group by
 - Specify the column(s) where you want to group (**key**)
 - Obtain a DataFrameGroupBy object

```
df = pd.DataFrame({'k' : ['a','b','a','b'],  
                  'c1': [2,10,3,15], 'c2' : [4,20,5,30]})  
grouped_df = df.groupby('k')    # 2 groups: 'a' and 'b'
```

Index	k	c1	c2
0	a	2	4
1	b	10	20
2	a	3	5
3	b	15	30



Index	k	c1	c2
0	a	2	4
2	a	3	5
1	b	10	20
3	b	15	30



■ Iterating on groups

- Each group is a subset of the original DataFrame

```
In [1]: for key, group_df in grouped_df:  
        print(key)  
        print(group_df)
```

Out[1]:

a

	k1	c1	c2
0	a	2	4
2	a	3	5



Index	k1	c1	c2
0	a	2	4
2	a	3	5

b

	k1	c1	c2
1	b	10	20
3	b	15	30



Index	k1	c1	c2
1	b	10	20
3	b	15	30



- **Aggregating** by group (min, max, mean, std)
 - The output is a DataFrame with the result of the aggregation for each group

In [1]: `grouped_df.mean() # Mean, separately for each group`

Out[1]:

k	c1	c2
a	2.5	4.5
b	12.5	25.0

Index	k1	c1	c2
0	a	2	4
2	a	3	5
Index	k1	c1	c2
1	b	10	20
3	b	15	30

The index of the result is the key of each group



- **Many Aggregations** with `.agg([min, max, ...])`
 - The output is a DataFrame with as many results as the number of aggregations x `n_columns`

In [1]: `grouped_df.agg(['max', 'min'])` # *max and min of each column for each group*

Out[1]:

		c1		c2		
k		max	min	max	min	
a		3	2	5	4	→
b		15	10	30	20	
						↑ ↑ ↑ ↑

Index	k1	c1	c2
0	a	2	4
2	a	3	5

Index	k1	c1	c2
1	b	10	20
3	b	15	30

The results now have 4 columns:
'max' and 'min' for each of the
previous columns



- **Aggregating** a single column by group
 - The output is a Series with the results of the aggregation for each group

```
In [1]: grouped_df['c1'].mean()
```

```
Out[1]:
```

```
k  
a    2.5  
b   12.5  
Name: c1, dtype=float64
```

Index	k1	c1	c2
0	a	2	4
2	a	3	5
Index	k1	c1	c2
1	b	10	20
3	b	15	30



■ Filtering data by group

- The filter is expressed with a lambda function working with each group DataFrame (x)

```
In [1]: # Keep groups for which column c1 has a mean > 5  
grouped_df.filter(lambda x: x['c1'].mean()>5)
```

Out[1]:

	k	c1	c2
1	b	10	20
3	b	15	30

Index	k1	c1	c2
0	a	2	4
2	a	3	5
Index	k1	c1	c2
1	b	10	20
3	b	15	30

mean = 2.5
x: filtered
out

mean = 12.5
x: kept in
the result



- Load DataFrame from **csv** file
 - Allows specifying the column **delimiter (sep)**
 - Automatically read **header** from first line of the file (after **skipping** the specified number of rows)
 - Column data types are inferred

```
df = pd.read_csv('./mycsv.csv', sep=',', skiprows=1)
```

mycsv.csv

MyTitle

c1,c2,c3

0,1,2

3,4,5

6,7,8



	c1	c2	c3
0	0	1	2
1	3	4	5
2	6	7	8



- Load DataFrame from **csv** file
 - If it contains **null** values, you can specify how to recognize them
 - Empty columns are converted to “NaN” (Not a Number)
 - Using `np.nan` (NumPy’s representation of NaN)
 - The string ‘NaN’ is automatically recognized

```
df = pd.read_csv('./mycsv.csv', sep=',',  
                 na_values=['no info', 'x'])
```

mycsv.csv

```
c1,c2,c3  
0,no info,  
3,4,5  
6,x,NaN
```



	c1	c2	c3
0	0	NaN	NaN
1	3	4.0	5.0
2	6	NaN	NaN

*type(np.nan) → float,
hence c2 and c3 are floats*



- **Save DataFrame to csv**

```
df.to_csv('./savedcsv.csv', sep=',')
```

	c1	c2	c3
0	0	NaN	2
1	3	4	5
2	6	NaN	NaN



savedcsv.csv

```
c1,c2,c3  
0,0,,2  
1,3,4,5  
2,6,,
```

- Use **index=False** to avoid writing the index

```
df.to_csv('./savedcsv.csv', sep=',', index=False)
```



- Load DataFrame from **json** file

```
df = pd.read_json('./myjson.json')
```

myjson.json

```
{"c1":{"0":0, "1":3, "2":6},  
 "c2":{"0":null, "1":4, "2":null},  
 "c3":{"0":2, "1":5, "2":null}}
```



	c1	c2	c3
0	0	NaN	2
1	3	4	5
2	6	NaN	NaN

- Use **pd.to_json(path)** to save a DataFrame in json format




DataFrames and I/O

- Many other data types are supported
 - Excel, HTML, HDF5, SAS, ...
- Check the pandas documentation
 - https://pandas.pydata.org/pandas-docs/stable/user_guide/io.html



Notebook Examples


- **3.2 Pandas
Grouping.ipynb**





Extra Notebook Examples

- **3.3 Pandas
Operations.ipynb
(Extra)**





- To know more about:
 - Missing values
 - Pivoting
 - Multi-indexing
- Check the set of slides «3-Pandas (extra).pdf»